

Rayfield-Based Calibration and Physical Model Identification of a Common Main Objective Stereo Microscope

Jean-François Witz*

Univ. Lille, CNRS, Centrale Lille, UMR 9013 – LaMcube

Laboratoire de Mécanique, Multiphysique, Multiéchelle, F-59000 Lille, France

May 24, 2026

Abstract

Common Main Objective (CMO) stereo microscopes violate the single-optical-centre assumption of standard camera calibration, making direct joint estimation of optical parameters and board poses poorly conditioned. A CMO-like direct bundle-adjustment diagnostic gives a Schur-complement coupling norm of 0.81, which is used here as a conditioning indicator rather than as an independent metrological measurement. We present a two-stage approach: first, a Zernike rayfield—a per-pixel 3D line—is fitted to ChArUco corner observations, decoupling measurement from model identification (coupling norm 0.0); second, a compact physical model is constructed by iterative residual analysis, where the Zernike decomposition of the residual between candidate model and measured rayfield reveals which physical degree of freedom is missing. On 10 stereo pairs from the open Pycaso dataset, the final 26-parameter model—a quasi-telecentric CMO skeleton with per-channel SE(3) arm alignment—achieves 1.06 px reprojection error (P50=0.87 px, P95=1.84 px), compared to > 300 px for standard OpenCV stereo calibration. All geometric descriptors (baseline, working distance, convergence angle) are reported within the chosen rayfield gauge and fixed reference focal scale ($f_x = 25600$ px); their absolute metrological values require an independent external reference. A ray-space BIC comparison followed by an operational reprojection guard selects this model as the best usable physical model; the guard is an engineering usability criterion, not a likelihood-derived BIC term. The Zernike rayfield serves not merely as a calibration model but as a diagnostic instrument: it tells the user *which* physical degree of freedom is missing, guiding hypothesis refinement through residual structure rather than blind fitting.

Keywords: stereo microscope; CMO; non-central calibration; Zernike rayfield; BIC model selection; telecentricity

1 Introduction

Stereo calibration is a mature technology when the camera obeys the pinhole model: all chief rays pass through a single optical centre, and mature toolchains (OpenCV, MATLAB, Bouguet) produce reliable results. This assumption breaks down for Common Main Objective (CMO) stereo microscopes, which use a single large objective shared between two off-axis sub-apertures [10, 22]. Each channel’s chief rays appear to originate from a different effective sub-pupil, separated by the stereo baseline. OpenCV’s

*ORCID: 0000-0002-7240-9476. Corresponding author: jean-francois.witz@centralelille.fr

stereoCalibrate [16] assumes one optical centre per camera; on the CMO architecture it produces unusable results (> 300 px reprojection error under the tested configuration).




Figure 1. CMO stereo microscope optical layout. Two off-axis sub-pupils share a single main objective. Chief rays converge toward the object plane through different effective origins—not through a single common centre.

The broader problem is that many optical instruments—CMO and Greenough stereo microscopes, endoscopes, Scheimpflug cameras, multi-camera arrays—depart from the central projection model. Generic camera models [9, 29, 24] and recent stereo microscope calibration studies [25, 34, 12] highlight the need for non-central approaches, but do not provide a systematic method for identifying the correct physical model family from observations (see Section 2 for a detailed review).

StereoComplex [32] addresses this gap: it fits a Zernike rayfield $\mathcal{R}(u, v) = (O(u, v), d(u, v))$ —a 3D line for every pixel—to corner observations, using a Zernike polynomial basis [33, 15]. This two-stage decomposition—measure rays first, identify optics second—is the central insight: it decouples the severely ill-conditioned joint estimation of optical parameters and board poses into a well-conditioned rayfield fit followed by pose-free model selection [24]. From this *measured* rayfield, physical descriptors can be read directly (Section 3), structural mismatches diagnosed by residual modal analysis (Section 3.3), and compact physical models constructed iteratively (Section 3.4–3.5).

This paper presents the first real-data validation of the StereoComplex rayfield-mediated non-central calibration workflow on a CMO stereo microscope using legacy ChArUco [8, 21] images from the open Pycaso dataset [7]. The contributions are:

1. A complete pipeline from raw ChArUco images to a compact 26-parameter physical CMO model,

Two-stage decomposition: measure rays first, identify optics second

Stage 1 stays in 2D (no 3D model assumption); Stage 2 identifies optics in ray space.




Figure 2. Pipeline overview. The two-stage decomposition separates rayfield measurement (Ray2D preprocessing + Zernike BA) from physical model identification (ray-space BIC selection). The Ray2D stage is purely 2D and does not impose a 3D camera model.

validated on real data.

2. A residual analysis method—Zernike modal decomposition of direction and moment differences between a candidate physical model and the measured rayfield—that identifies missing degrees of freedom.
3. A two-level model-selection framework: ray-space BIC is used to identify the optical family, while a separate operational reprojection guard rejects models whose pixel RMS is not usable for calibration.
4. Demonstration that the Zernike rayfield is an excellent initialiser for direct corner bundle adjustment: a joint model+pose refinement (200 nfev, soft-L1 loss) improves pixel RMS from 1.06 px to 0.88 px (17 %) while leaving all 26 model parameters unchanged to within 10^{-3} , confirming that the rayfield-initialised parameters are already near-optimal for corner reprojection.

2 Related Work

2.1 Generic and Non-Central Camera Calibration

The pinhole camera model with polynomial distortion [35, 5] remains the standard in computer vision toolkits [16, 4]. For wide-angle and fisheye lenses, unified models [13, 11, 23] extend the pinhole paradigm with additional radial terms. However, these parametric models share a fundamental assumption: all chief rays pass through a single optical centre.

For cameras that violate this assumption, generic (non-parametric) models assign a per-pixel ray without imposing a central projection constraint. Grossberg and Nayar [9] introduced the raxel model, treating each pixel as an independent ray. Sturm and Ramalingam [29], later refined by Ramalingam and Sturm [19], developed a generic calibration framework from multiple views of a planar target, showing that a flexible pixel-to-ray mapping can be estimated without knowing the camera model *a priori*. Miraldo and Araujo [14] proposed smooth generic camera models using continuous ray-field interpolation. Soloff et al. [27] used a polynomial expansion for 3D particle image velocimetry. Camposeco et al. [6] demonstrated spline-based generic calibration with subpixel accuracy. Ramalingam, Sturm, and Nayar [20] unified these approaches in a generic structure-from-motion framework.

Most recently, Schöps et al. [24] showed that highly parameterised generic models (10,000+ parameters) can outperform classical 12-parameter models on standard benchmarks, challenging the assumption that parametric models are inherently more robust.

Our work differs from these generic approaches in two ways: (1) we use the Zernike rayfield not as a final calibration model but as a *diagnostic instrument* to discover the correct physical model family, and (2) we construct a compact, interpretable physical model guided by residual structure rather than imposing a generic non-parametric fit.

2.2 Stereo Microscope Calibration

Stereo microscopes pose a specific calibration challenge: the two optical channels share a common main objective but have separate effective pupils [25]. Digital image correlation under optical microscopes requires careful distortion correction [34, 12].

These methods calibrate a specific instrument but do not address the generic problem of *identifying* the correct optical model family from observations. Our contribution is a systematic framework that

measures the ray-field first, then uses residual analysis to construct a physical model that matches the observed structure—here, discovering the telecentric CMO family and the SE(3) arm misalignment.

2.3 Model Selection in Camera Calibration

Model selection in calibration traditionally relies on reprojection RMS, which penalises underfitting but is insensitive to overfitting. The Akaike Information Criterion (AIC) [1] and Bayesian Information Criterion (BIC) [26] add parameter-count penalties, providing a principled trade-off between fit quality and model complexity. Strecha et al. [28] established benchmark datasets and evaluation protocols for camera calibration and multi-view stereo. Our work extends BIC-based model selection to ray space: we compare optical model families (pinhole, Brown-Conrady, parallel-plate, telecentric CMO) against a measured Zernike rayfield, and add an operational reprojection guard that rejects models whose pixel RMS exceeds a usability threshold—a constraint that the raw ray-space BIC does not capture.

2.4 Plücker Coordinates and Line-Based Geometry

Plücker coordinates [18] represent a 3D line as a 6-vector (d, m) where d is the direction and $m = O \times d$ is the moment. They are widely used in geometric computer vision for line-based structure from motion and calibration [2, 30]. Our residual analysis uses Plücker moments to separate direction errors from origin errors: Δd reveals angular misalignment; Δm reveals the effective ray origin shift. The Z_0^0 -dominated residual in both Δd and Δm is the key diagnostic that guides the SE(3) arm alignment.

2.5 Residual-Guided Model Construction

The idea of using residuals to guide model refinement has roots in adaptive optics and wavefront sensing [33, 15], where Zernike decomposition of a distorted wavefront identifies specific aberrations (defocus, astigmatism, coma) that are then corrected by deformable mirrors. Our method applies the same principle to camera calibration: the Zernike decomposition of the residual between a candidate physical model and the measured rayfield identifies which physical degree of freedom is missing—a Z_0^0 -dominated residual points to a global offset (SE(3) arm misalignment); a Z_2 -dominated residual would point to astigmatism or field curvature. This is distinct from generic non-parametric approaches [24] that fit a flexible model without extracting physical meaning from the residual structure.

3 Method

3.1 Zernike Order Selection by BIC

The Zernike rayfield has two hyperparameters: the origin polynomial order o and the direction polynomial order d . Increasing either adds Zernike modes, reducing pixel RMS at the cost of more parameters. We sweep $o \in \{0, 1, 2\}$ and $d \in \{2, 3, 4\}$ (7 combinations) and select the optimal order by BIC [26]:

$$\text{BIC}_{\text{px}} = k \log(N_{\text{obs}}) + N_{\text{obs}} \log(\text{RSS}_{\text{px}}/N_{\text{obs}}), \quad (1)$$

where k is the total number of parameters (Zernike coefficients + poses), $N_{\text{obs}} = 3300$ corner observations, and RSS_{px} is the sum of squared pixel-equivalent residuals.




Figure 3. Zernike order selection by BIC. Pixel RMS decreases monotonically with more parameters (from 0.47 px at 57 params to 0.41 px at 141 params), but BIC penalises complexity. The BIC minimum is at O(1)+d(3) (93 params, BIC = -5093 , RMS = 0.41 px), though O(0)+d(2) (57 params, BIC = -4522) is within $\sim 10\%$ in RMS at 40 % fewer parameters. The O(0)+d(2) model (57 params, 0.47 px) is selected for physical model construction because it is internally consistent with the rigid-sub-pupil CMO assumption, achieves 90 % of the noise-floor accuracy at nearly half the parameters of the BIC minimum, and the marginal improvement above d(2) falls within corner detection noise.

Table 1. Zernike order sweep: pixel RMS and BIC.

Model	k	RMS (px)	BIC _{px}
O(0)+d(2)	57	0.470	−4522
O(1)+d(2)	69	0.444	−4803
O(0)+d(3)	81	0.419	−5080
O(1)+d(3)	93	0.412	−5093
O(2)+d(3)	111	0.409	−4995
O(1)+d(4)	123	0.412	−4858
O(2)+d(4)	141	0.410	−4743

Table 1 summarises the full sweep. The BIC minimum is therefore interpreted as a diagnostic of the best flexible rayfield accuracy available in this data regime. The O(0)+d(2) rayfield is not claimed to be the statistical BIC optimum; it is deliberately retained as the physically consistent reference for CMO model construction because it enforces rigid effective sub-pupils while remaining within 0.06 px of the BIC-minimum RMS.

The BIC minimum at O(1)+d(3) achieves 0.41 px RMS with 93 parameters. We nevertheless select O(0)+d(2) (57 params, 0.47 px) as the preferred rayfield configuration, for three reasons:

1. **Simplicity:** 57 params vs. 93 — nearly half the parameters for a penalty of only 0.06 px in RMS (~ 0.1 px in practice).
2. **Marginal improvement:** the residual gain from d(2) to d(3) (0.47 px $\rightarrow$ 0.41 px) is well within typical subpixel noise of ChArUco detection; the additional BIC improvement does not justify the extra model complexity.
3. **Consistency with the physical model:** the CMO physical model assumes rigid sub-pupils ($o = 0$). Using the same origin order in the Zernike rayfield makes the two representations internally consistent and comparable: the Zernike field measures what the physical model can represent.

Selecting O(0)+d(2) achieves 90 % of the noise-floor accuracy with a representation that is internally consistent with the physical model it validates.

3.2 Per-Pixel Ray-Field and Transverse Gauge

For each pixel (u, v) on the sensor, the imaging system defines a 3D line in space. We represent this line $\mathcal{L}_c(u, v)$ by an origin point $O_c(u, v) \in \mathbb{R}^3$ and a unit direction vector $d_c(u, v) \in \mathbb{S}^2$, where $c \in \{L, R\}$ denotes the left or right channel. Points on the line are parameterised as $O_c + t d_c$ for $t \in \mathbb{R}$.

The origin $O_c(u, v)$ is not unique: any point along the ray would represent the same line. We fix this gauge freedom by enforcing the **transverse gauge** $O_c \cdot d_c = 0$, which selects the point on the ray closest to the coordinate origin. This eliminates one redundant degree of freedom per ray and stabilises the parameterisation.

The origin and direction fields are expanded on a Zernike polynomial basis [33, 15]:

$$O_c(u, v) = \sum_m a_{c,m}^O Z_m(\tilde{u}, \tilde{v}), \quad d_c(u, v) = \text{normalize} \left(d_c^{\text{pinhole}}(u, v) + \sum_{m \leq d_{\max}} a_{c,m}^d Z_m(\tilde{u}, \tilde{v}) \right), \quad (2)$$

where Z_m are the real-valued Zernike polynomials on the unit disk, $\tilde{u}, \tilde{v}$ are normalised pixel coordinates mapping the sensor to $[-1, 1]^2$, d_c^{pinhole} is the initial pinhole direction from a fixed central projection matrix K , and the direction perturbation is projected transverse to d_c^{pinhole} before renormalisation. The expansion order for the origin field is $o_{\max}$ and for the direction field is $d_{\max}$; selecting $o_{\max} = 0$ and $d_{\max} = 2$ (see Section 2.1) gives 3 origin coefficients and 18 direction coefficients per channel.

3.3 Ray-Space Bundle Adjustment

The rayfield coefficients $a_{c,m}^O, a_{c,m}^d$ and the board poses $\{\eta_p\}_{p=1}^P$ are estimated jointly from N ChArUco corner observations. Each observation k links a pixel (u_k, v_k) on channel c_k to a known 3D board point X_k expressed in the board frame. The board pose $\eta_p = (R_p, t_p) \in \text{SE}(3)$ transforms X_k to the camera frame: $X_k^{\text{cam}} = R_p X_k + t_p$.

The **ray-space residual** measures the perpendicular distance from the transformed board point to the ray $\mathcal{L}_{c_k}(u_k, v_k)$:

$$r_k = \|(X_k^{\text{cam}} - O_{c_k}) - ((X_k^{\text{cam}} - O_{c_k}) \cdot d_{c_k}) d_{c_k}\|. \quad (3)$$

This is the norm of the component of $X_k^{\text{cam}} - O_{c_k}$ orthogonal to the ray direction d_{c_k} . The objective is the sum of squared residuals:

$$\min_{\{a_{c,m}\}, \{\eta_p\}} \sum_k r_k^2, \quad (4)$$

solved by nonlinear least-squares.

For the CMO microscope, we impose the physically-motivated constraint that the board is mounted on a Z-only translation stage: all P frames share a common rotation R and XY translation (t_x, t_y) , with only the Z component varying per frame. This reduces the pose parameters from $6P$ to $3 + 2 + P = 15$ (for $P = 10$ frames: $60 \rightarrow 15$), making the 57-parameter Zernike BA well-conditioned.

The ray-space BA differs from classical reprojection minimisation [35] in a fundamental way: it optimises the 3D geometric consistency of rays rather than the 2D reprojection error. This is the measurement stage of the two-stage decomposition (see Section 4.2).

3.4 The Zernike Rayfield as Observable

The fitted rayfield $\mathcal{R}_c(u, v) = (O_c(u, v), d_c(u, v))$ is used as an experimental observable. Within the chosen transverse gauge and fixed reference scale, effective rayfield descriptors can be read from the centre-pixel rays—no physical model fit is required. These descriptors are internally consistent quantities of the calibrated rayfield; they are not claimed to be externally validated absolute microscope specifications without an independent metrological reference. The centre-pixel values are: $b = \|O_R - O_L\| = 24.9$ mm (stereo baseline), $z_p = (|O_{L,z}| + |O_{R,z}|)/2 = 2.5$ mm (sub-pupil depth), $WD = 64.7$ mm (working distance, from mean pose Z), $f_{\text{obj}} = WD - z_p = 62.2$ mm (objective focal length), $\theta = \arccos(d_L \cdot d_R) = 22.6^\circ$ (convergence angle).

3.5 Physical Model Construction by Residual Analysis

The core methodology is iterative: propose a candidate physical model, compute its residual against the Zernike rayfield, decompose the residual on Zernike modes, and use the dominant mode to identify the missing degree of freedom.




Figure 4. 3D reconstruction of the effective sub-pupils from the Zernike rayfield. O_L and O_R are read from the centre-pixel rayfield. The chief rays converge at the working plane ($WD = 64.7$ mm). The baseline $b = \|O_R - O_L\| = 24.9$ mm is an effective rayfield descriptor in the selected gauge and reference scale.

3.5.1 Perspective CMO (Baseline)

The simplest CMO model places a perspective camera behind each decentred sub-pupil $S_c = (\pm b/2, 0, WD - f_{\text{obj}})$. This predicts $d_y(u, v) \propto (v - c_y)$ —a strong linear gradient. The measured d_y field is nearly constant across the sensor (range 0.079, mean +0.059), while the perspective model predicts a range of 0.232—a $3\times$ discrepancy. The real system is **object-space telecentric**, not perspective.




Figure 5. Comparison of $d_y(u, v)$ profiles across the sensor. The Zernike rayfield (measured) shows a nearly constant d_y (range 0.079), while the perspective CMO model predicts a $3\times$ larger linear gradient (range 0.232). The telecentric model reproduces the measured profile almost exactly.

3.5.2 Telecentric CMO

The perspective model fails because the real microscope is object-space telecentric: chief rays are nearly parallel rather than diverging from a point. The telecentric model directly parameterises the ray-field without deriving it from a point projection. For each channel $c \in \{L, R\}$, the origin is a rigid sub-pupil:

$$O_c = S_c = (\pm b/2, 0, WD - f_{\text{obj}}), \quad (5)$$

where the sign is $+$ for right and $-$ for left. The direction field is an affine function of normalised angular coordinates $\tilde{u} = (u - c_x)p/f_{\text{ang}}$, $\tilde{v} = (v - c_y)p/f_{\text{ang}}$, where p is the pixel pitch, f_{ang} is the angular focal length, and (c_x, c_y) is the principal point:

$$d_c(u, v) = \text{normalize}(d_{c,0} + s_{x,c} \tilde{u} e_x + s_{y,c} \tilde{v} e_y + s_{xy,c} \tilde{u} \tilde{v} e_x + s_{yx,c} \tilde{v} \tilde{u} e_y), \quad (6)$$

where $d_{c,0} = (\pm \sin \theta, d_y^{\text{com}}, \cos \theta)$ is the chief-ray direction, with θ the convergence half-angle and d_y^{com} a shared vertical component.

Adding pupil shear—an affine transverse variation of the origin—completes the model:

$$O_c(u, v) = S_c + P_{d_c}^\perp(\rho_{x,c} \tilde{u} e_x + \rho_{y,c} \tilde{v} e_y), \quad (7)$$

where $P_d^\perp v = v - (v \cdot d) d$ projects onto the plane orthogonal to d , ensuring $O_c \cdot d_c = 0$.

The 14 parameters are: f_{obj} , WD , b (skeleton), c_x , c_y (principal point), f_{ang} (angular scale), θ , d_y^{com} (chief-ray direction), $s_{x,L}$, $s_{y,L}$, $s_{x,R}$, $s_{y,R}$ (per-channel slopes), ρ_x , ρ_y (shared pupil shear). This model achieves 0.12 mm ray-space RMS (down from 3.48 mm for perspective) and 14.6 px reprojection (down from 86 px). The dominant geometry is captured, but 14.6 px is far from the Zernike reference (0.47 px).

3.5.3 Residual Analysis

The residual between the telecentric model and the Zernike rayfield is computed on a 41×41 grid: $\Delta d = d_{\text{Zernike}} - d_{\text{CMO}}$, $\Delta m = m_{\text{Zernike}} - m_{\text{CMO}}$, where $m = O \times d$ is the Plücker moment [18].

Projecting on Zernike modes up to order 4 reveals that **both Δd and Δm are 97–98 % Z_0^0 (global piston)**. This is the crucial diagnostic: the residual is not a spatial field distortion (which would appear in higher Zernike modes) but a **global line-bundle offset**—the entire set of rays from each channel is slightly rotated or translated relative to the Zernike reference.

3.5.4 Testing Alternative Hypotheses

Two alternative hypotheses are tested and rejected:

- **Image-space pre-warp:** a polynomial mapping $\xi = W(u, v)$ before the direction model. Affine and quadratic warps *degrade* pixel RMS ($14.6 \rightarrow 16.0\text{--}16.5$ px).
- **Spatially varying origin:** affine and quadratic transverse origin fields with direction fixed. Only 8 % ray RMS improvement—the residual is not spatial.

Both negative results are consistent with the Z_0^0 diagnostic (they would produce spatial, not global, changes).

3.5.5 Residual-Guided Hypothesis Construction

The central methodological contribution is that physical model construction is guided by *residual structure*, not blind fitting. At each stage, the direction residual $\Delta d = d_{\text{Zernike}} - d_{\text{model}}$ is computed on a 31×31 grid and displayed as a heatmap. The pattern of the residual directly reveals which physical degree of freedom is missing.

Stage 1 (Perspective CMO): the residual is large and structured, with RMS 23.6° (max 27.7°). The d_y component shows a strong linear gradient trace—perspective projection from a point sub-pupil—that the Zernike field does not have. This **rejects** the perspective family and points to object-space telecentricity.

Stage 2 (Telecentric CMO): the residual drops to 0.27° RMS (two orders of magnitude), confirming that the telecentric model family is correct. However, the residual is Z_0^0 -dominated (97–98 % global piston in both Δd and Δm)—a constant offset across the entire field. This **reveals** that the missing degree of freedom is a global line-bundle misalignment, not a spatial field distortion.

Stage 3 (CMO + SE(3)): the residual drops further to 0.003° RMS, near the noise floor. The remaining gap to the Zernike reference (0.47 px vs 1.06 px) comes from distributed low-amplitude aberrations (field curvature, astigmatism) that require the flexibility of the 57-parameter Zernike model.




Figure 6. Residual evolution across model stages. Direction error Δd (degrees) between the Zernike rayfield (reference) and each candidate model, evaluated on a 31×31 pixel grid. **Left:** Perspective CMO — large structured error (RMS 23.6° , max 27.7°), dominated by a strong d_y gradient. **Centre:** Telecentric CMO — RMS drops to 0.27° (two orders of magnitude), but the residual is still Z_0^0 -dominated (global piston), revealing the missing SE(3) arm alignment. **Right:** CMO + SE(3) — RMS drops to 0.003° , near the noise floor. The residual heatmap at each stage directly tells the user *what* is wrong and *where* to look next.

The critical point is that **at no stage was a model added blindly**. The heatmap pattern at each stage dictates the next hypothesis: perspective $\rightarrow$ telecentric (because the d_y gradient is wrong), telecentric $\rightarrow$ SE(3) (because the residual is Z_0^0 piston, not spatial).

3.5.6 Per-Channel SE(3) Arm Alignment

The Z_0^0 -dominated residual points to a global misalignment of each optical arm. A rigid transform $(R_c, t_c) \in \text{SE}(3)$ is applied to each channel’s ray bundle, representing a small rotation and translation of the entire optical arm relative to the ideal CMO skeleton. For each pixel (u, v) on channel c , the telecentric ray $(O_c^{\text{tel}}, d_c^{\text{tel}})$ is transformed:

$$d_c(u, v) = R_c d_c^{\text{tel}}(u, v), \quad O_c(u, v) = R_c O_c^{\text{tel}}(u, v) + t_c. \quad (8)$$

Each $R_c \in \text{SO}(3)$ is parameterised by a 3-element rotation vector $r_c \in \mathbb{R}^3$ ($\theta_c = \|r_c\|$, axis $r_c/\|r_c\|$). The 12 additional parameters (r_L, t_L, r_R, t_R) bring the total to 26.

Jointly fitting the 14 telecentric parameters and the 12 SE(3) parameters against the Zernike rayfield reduces pixel RMS from 14.6 px to 1.06 px (P50=0.87 px, P95=1.84 px)—a $14\times$ improvement. The fitted rotations are $\sim 2.5^\circ$ (left) and $\sim 3.7^\circ$ (right); translations are sub-mm. The rotation angles are stable across optimisation runs; the translation components trade off with the telecentric base parameters (a known identifiability limitation).

3.6 Model Selection

Ray-space BIC [26] identifies the correct optical family among five candidates (pinhole, Brown-Conrady, parallel-plate, telecentric CMO with/without shear). The telecentric CMO+shear model wins by $> 40\,000$ BIC points over all alternatives.

However, this 14-parameter model produces 14.6 px reprojection—it is the correct *family* but not a

usable model. We introduce an **operational usability score** that adds a hard reprojection guard:

$$S_{\text{usable}} = \text{BIC}_{\text{ray}} + \mathbf{1}_{e_{\text{px}} > 1.5 \text{ px}} \left[10^6 + N_{\text{px}} \log \left(\frac{e_{\text{px}}^2}{1.5^2} \right) \right]. \quad (9)$$

Figure 7 shows the gauge-regularisation sweep. The 14p telecentric model is rejected ($S_{\text{usable}} \approx +978\,890$); the 26p SE(3)-aligned model is selected as best usable ($S_{\text{usable}} \approx -32\,433$).




Figure 7. Pareto frontier of the gauge-regularized Zernike sweep. Pixel RMS vs. Z_0 direction drift for varying regularization strengths $\sigma_{Z_0}, \sigma_{Z_1}$. The near-vertical Pareto curve confirms that after double TPS preprocessing, even the unregularized full-pose fit has negligible gauge drift (0.023°). Regularization is unnecessary when 2D corners are sufficiently clean.

3.7 Why Direct Bundle Adjustment is Ill-Conditioned for Non-Central Optics

The two-stage decomposition strategy (rayfield measurement $\rightarrow$ model identification) is motivated by a fundamental conditioning problem in direct bundle adjustment (BA). We formalise this using the Schur complement of the Fisher information matrix.

3.7.1 Conditioning of the Joint Estimation Problem

Direct BA jointly estimates optical parameters $\theta \in \mathbb{R}^p$ and board poses $\eta \in \mathbb{R}^q$ ($q = 6P$ for P frames) by minimising the pixel reprojection error. The Fisher information matrix at the optimum has the block structure:

$$\mathcal{I}(\theta, \eta) = \begin{bmatrix} \mathcal{I}_{\theta\theta} & \mathcal{I}_{\theta\eta} \\ \mathcal{I}_{\eta\theta} & \mathcal{I}_{\eta\eta} \end{bmatrix} = J^T J, \quad (10)$$

where $J = \partial r / \partial(\theta, \eta)$ is the residual Jacobian. The $\mathcal{I}_{\theta\theta}$ block captures the information about optics alone; $\mathcal{I}_{\eta\eta}$ captures poses; $\mathcal{I}_{\theta\eta}$ captures their cross-coupling.

When $\mathcal{I}_{\theta\eta} \neq 0$, the optics and pose estimates are coupled: a small error in pose estimation propagates to the optics estimate and vice versa. The effective information about θ after marginalising over the uncertain poses is the **Schur complement**:

$$\mathcal{I}_{\theta|\eta} = \mathcal{I}_{\theta\theta} - \mathcal{I}_{\theta\eta} \mathcal{I}_{\eta\eta}^{-1} \mathcal{I}_{\eta\theta}. \quad (11)$$

We define the **coupling norm** $c \in [0, 1]$ as the fraction of the optics information that is lost to pose coupling:

$$c = \frac{\|\mathcal{I}_{\theta\eta} \mathcal{I}_{\eta\eta}^{-1} \mathcal{I}_{\eta\theta}\|_F}{\|\mathcal{I}_{\theta\theta}\|_F}, \quad (12)$$

where $\|\cdot\|_F$ is the Frobenius norm. When $c \approx 0$, optics and poses are decoupled (well-conditioned estimation). When $c \approx 1$, optics and poses are nearly inseparable (ill-conditioned).

3.7.2 Empirical Coupling Values

The Zernike rayfield itself (57 parameters, 0.47 px RMS) serves as a generic non-parametric reference: it is a per-pixel spline in the Zernike basis and makes no assumptions about the optical architecture. The compact 26p model loses approximately 0.6 px relative to this ceiling. This 0.6 px gap represents the irreducible cost of physical interpretability: the 26 parameters capture the dominant telecentric skeleton and SE(3) arm misalignment, while the remaining 0.6 px come from low-amplitude aberrations (field curvature, astigmatism, detection noise) that only the flexible 57-parameter Zernike field can absorb. For applications where sub-pixel accuracy is less critical than having a physically meaningful model, this trade-off is acceptable.

In the StereoComplex direct-BA diagnostic implementation, a CMO-like configuration yields a Schur-complement coupling norm $c = 0.81$, computed from a finite-difference Jacobian (forward differences, step $\delta = 10^{-6}$, parameters scaled to comparable units). Because such block norms depend on the parameter scaling, the value is reported as a conditioning diagnostic rather than as a dimensionless property of the microscope. This value is reported as a conditioning diagnostic of the inverse problem, not as an independent physical measurement on the Pycaso dataset. Because such block norms can depend on parameter scaling, the value should be interpreted qualitatively: it indicates strong optics–pose coupling, not a universal dimensionless property of all CMO microscopes. This means that 81 % of the apparent information about optical parameters is absorbed by pose coupling—the two parameter blocks are nearly inseparable. By contrast, the rayfield-mediated pipeline operates with $c = 0.0$: the Zernike BA absorbs pose uncertainty into the rayfield through its constrained shared-rotation model (15 pose parameters), and the ray-space model identification stage is entirely pose-free.

Figure 8 shows the singular value spectra. The coupling norm $c = 0.81$ quantifies why direct BA fails to converge on the Pycaso CMO oracle (RMS 29 px, Section 4): the optimiser cannot separate optical parameters from pose parameters, and the pinhole-based initialisation (≈ 2 px error) is too far from the true solution. Schur-coupling values for the other oracles (pinhole, Brown-Conrady, parallel-plate, Greenough) were not reliably computable from the current diagnostic infrastructure and are therefore not reported.

3.8 Corner Bundle Adjustment

The 26p model, fitted to the rayfield, is refined by direct corner bundle adjustment (joint optimisation of model parameters and per-frame poses on 3300 corner observations, 200 nfev, soft-L1 loss). The pixel RMS improves from 1.06 px to 0.88 px (P50=0.64 px, P95=1.70 px; 17 %), while all 26 model parameters remain unchanged to within 10^{-3} of their rayfield-fit values. The model geometry—baseline, working distance, focal length, convergence angle, SE(3) arm corrections—is essentially identical before and after corner refinement. This confirms that the rayfield-initialised parameters are already near-optimal for corner reprojection: the corner BA improves pose estimates but does not alter the physical model. This is a strong validation of the two-stage decomposition strategy: the rayfield fit, which never directly minimises corner error, produces parameters so close to the corner optimum that the dedicated BA can barely improve them. In contrast, the direct approach—jointly estimating optics and poses from




Figure 8. Singular value spectra of the optics information matrix $\mathcal{I}_{\theta\theta}$ (direct BA), its Schur complement $\mathcal{I}_{\theta|\eta}$ after marginalising poses, and the rayfield BA (pose-free). The trailing singular values of $\mathcal{I}_{\theta\theta}$ collapse in the Schur complement, indicating that the last ~ 8 optical parameters are nearly unobservable when poses are jointly estimated. The rayfield BA eliminates this coupling by design.

corners—is severely ill-conditioned for this non-central system (Schur-complement coupling norm 0.81), failing to converge with a pinhole initialisation (RMS 29 px).

3.9 Stabilising Direct BA with a Schur-Complement Prior

The corner BA of Section 3.8 improves pixel RMS by 17% without altering the physical model. However, the Schur-complement diagnostic of Section 3.7 reveals that the Fisher information matrix at the rayfield solution θ_0 has weak eigenmodes—directions in parameter space where a small change in optics can be nearly perfectly compensated by a change in board poses. An unregularised BA can drift along these modes, trading physical interpretability for marginal reprojection gains.

3.9.1 Schur-Eigenmode Prior

The Schur complement at the rayfield solution is diagonalised,

$$S_0 = S_\theta(\theta_0, \eta_0) = V\Lambda V^T, \quad \Lambda = \text{diag}(\lambda_1, \dots, \lambda_{n_\theta}), \quad \lambda_1 \geq \dots \geq \lambda_{n_\theta}, \quad (13)$$

where $V = [v_1, \dots, v_{n_\theta}]$ are the eigenvectors and λ_i the eigenvalues. Directions v_i with small $\lambda_i/\lambda_{\max}$ are poorly observable after pose marginalisation.

The prior penalises displacement along each eigenmode with a weight inversely proportional to its eigenvalue [31]:

$$\mathcal{L}_{\text{Schur}}(\theta) = \alpha \sum_{i=1}^{n_\theta} w_i (v_i^T D_\theta^{-1} (\theta - \theta_0))^2, \quad w_i = \left(\frac{\lambda_{\max}}{\lambda_i + \varepsilon \lambda_{\max}} \right)^p, \quad (14)$$

where $D_\theta = \text{diag}(\sigma_1, \dots, \sigma_{n_\theta})$ contains per-parameter scales (degrees for rotations, millimetres for translations, pixels for the principal point), α controls the overall prior strength, p the sharpness of the eigenvalue weighting ($p = 1$ throughout), and $\varepsilon = 10^{-6}$ prevents division by zero. The prior is added to the least-squares objective as n_θ extra residual rows $r_i^{\text{prior}} = \sqrt{\alpha w_i} v_i^T D_\theta^{-1} (\theta - \theta_0)$.

3.9.2 Prior Sweep and Comparison with Tikhonov Baseline

An isotropic (Tikhonov) prior $\mathcal{L}_{\text{iso}} = \alpha \|D_\theta^{-1}(\theta - \theta_0)\|^2$ serves as a baseline. Both priors are swept over $\alpha \in \{10^{-4}, 10^{-3}, 10^{-2}, 10^{-1}, 1, 10\}$ with the same BA settings as Section 3.8 (200 nfev, soft-L1 loss).

Table 2. Regularisation sweep: isotropic (Tikhonov) vs. Schur-complement prior. Δ_{weak} is the norm of the scaled parameter displacement projected onto the weak Schur eigenspace ($\lambda_i/\lambda_{\text{max}} < 10^{-3}$). $\checkmark/\times$ indicates convergence.

Isotropic prior			Schur prior		
α	RMS (px)	Δ_{weak}	α	RMS (px)	Δ_{weak}
10^{-4}	0.239 ($\times$)	0.529	10^{-4}	0.277 ($\checkmark$)	0.0033
10^{-3}	0.245	0.248	10^{-3}	0.278	0.0007
10^{-2}	0.257	0.082	10^{-2}	0.279	0.0001
10^{-1}	0.270	0.017	10^{-1}	0.279	$< 10^{-4}$
10^0	0.278	0.003	10^0	0.283	$< 10^{-4}$
10^1	0.286	0.010	10^1	0.323	$< 10^{-4}$

Table 2 reports the results. The isotropic prior faces a classical trade-off: a small α leaves the weak modes uncontrolled ($\Delta_{\text{weak}} = 0.53$ at $\alpha=10^{-4}$, where the optimiser also fails to converge), while a large α degrades the fit (RMS rises to 0.286 px). The Schur prior breaks this trade-off: even at $\alpha=10^{-4}$ it suppresses 99.4% of the weak-mode drift while keeping the RMS within 0.039 px of the unregularised baseline. At $\alpha=10^{-3}$ the weak-mode drift is below 10^{-3} and the RMS penalty is only 0.039 px—the prior blocks the poorly observable directions while leaving the well-observed degrees of freedom free to improve the fit. At the strongest setting ($\alpha=10$), the Schur prior still converges, while the isotropic prior produces a higher RMS (0.286 px vs 0.323 px) with residual weak drift (0.010).

3.9.3 The Double Role of the Rayfield

This result formalises a double role for the rayfield estimate θ_0 in direct bundle adjustment:

1. **Initialiser.** θ_0 places the direct BA in the correct convergence basin (Section 3.8), avoiding the local minima that trap a pinhole or perspective-CMO initialisation (RMS 29 px vs. 0.88 px).
2. **Observability prior.** The Schur eigenmodes of the Fisher matrix at θ_0 tell the optimiser *which directions it may trust*. The prior blocks compensation between poses and intrinsics without penalising the genuinely observable optical degrees of freedom.

The Schur prior is not a generic regulariser—it is *derived from the same data* that produced the initialisation. The rayfield’s Fisher information matrix, evaluated at θ_0 , encodes the local observability structure of the inverse problem. Using that structure to guide the subsequent BA is the methodological contribution: the rayfield is not merely a calibration model but an *observability oracle* for direct refinement.

4 Results

4.1 Data Sources

Table 3 clarifies which results come from real Pycaso calibration data, which from synthetic oracles, and which are methodological diagnostics.

Table 3. Data sources for results reported in this paper. Real data: 10 Pycaso stereo pairs. Synthetic oracles: StereoComplex model_selection_oracles.py. Method: analytical and numerical diagnostics independent of data source.

Result category	Source	Examples
Real data	Pycaso CMO images (10 stereo pairs)	Zernike rayfield fit (0.47 px), CMO+SE(3) model (1.06 px), geometric descriptors (b , WD , θ), pose sweep, cross-validation, bootstrap, fx sensitivity
Synthetic oracles	StereoComplex simulator (noiseless)	Direct vs. rayfield inversion (coupling $c = 0.81$), ray-space BIC model selection, oracle parameter recovery
Methodological	Analytical / numerical diagnostics	BIC formula, Plücker residual decomposition, SE(3) ablation, Schur complement derivation

4.2 Dataset and Pipeline

Ten stereo pairs of 2048×2048 px ChArUco [8, 21] images (16×12 squares, 0.3 mm, DICT_6X6_250, setLegacyPattern(True)) from the Pycaso open dataset [7] are used. The Z-range is 2.65 mm–3.35 mm ($\Delta = 0.70$ mm). Corner detection yields 61–161 corners per frame; Hessian-based completion ($|\det H| + \text{Otsu} + \text{barycentre}$) fills all 165. A double Ray2D TPS denoising pass [3] (ArUco markers at $\lambda = 10$, then completed 165 corners at $\lambda = 3$, Huber $c = 1.5$) produces gauge-stable 2D corners. The two-phase TPS is not cosmetic: without it, raw ChArUco detections carry structured biases from compression artifacts, blur, and defocus that cause the Zernike rayfield fit to become gauge-unstable—the direction piston Z_0 drifts by 8.5° between constrained and full-pose fits, making physical interpretation impossible. After double TPS, the gauge drift drops to 0.023° , and the rayfield becomes a stable experimental oracle [32].

4.3 Pose Count Sensitivity

To assess robustness to the number of calibration frames, we sub-sample $N \in \{3, 5, 8, 10\}$ frames from the 10 available and re-run the full pipeline (Zernike fit + 26p model) for each. Table 4 reports pixel RMS and geometric descriptors.

Table 4. Pose sweep: pixel RMS and CMO descriptor stability vs. number of frames.

N frames	RMS (px)	P50 (px)	b (mm)	WD (mm)	θ ($^\circ$)
3	1.02	0.84	24.7	64.6	22.4
5	1.03	0.84	24.8	64.6	22.5
8	1.06	0.87	24.9	64.7	22.6
10	1.07	0.88	24.8	64.7	22.4

The pixel RMS is stable within 0.05 px across frame counts—even 3 frames suffice for 1.02 px calibration. Geometric descriptors vary by less than 0.2 mm (baseline) and 0.1 mm (working distance),

confirming that the pipeline is robust to the number of calibration images. The convergence angle $\theta \approx 22.5^\circ$ (standard deviation 0.1°).

4.4 OpenCV Tuning Does Not Recover

We tested the full OpenCV stereo calibration flag space: `CALIB_RATIONAL_MODEL` (8 distortion coefficients), `CALIB_THIN_PRISM`, and `CALIB_TILTED_MODEL`. None of these configurations produced a usable calibration on the CMO data. The rational model, with the most parameters (14 intrinsics per camera), still fails because the distortion model remains radial and centred—it cannot represent a non-central projection where rays originate from two different sub-pupils. The best OpenCV configuration achieves >300 px RMS, compared to 1.06 px for our 26-parameter CMO+SE(3) model. This confirms that the failure is structural (central projection assumption), not a matter of under-parameterised distortion.

4.5 Geometric Descriptor Validation

The Pycaso software publication [7] documents a Leica 205C stereo microscope. However, nominal geometric parameters (baseline, working distance, focal length) of the specific instrument used to acquire the open dataset are not available from the publication or manufacturer datasheets. Our rayfield-derived descriptors— $b = 24.9$ mm, $WD = 64.7$ mm, $f_{\text{obj}} = 62.2$ mm—are substantially different. This indicates that the open Pycaso calibration dataset was produced with a different microscope or optical configuration than the one documented in the Pycaso paper.

While we cannot validate against the nominal specifications of the actual microscope used (unknown manufacturer, unknown model), the StereoComplex-derived descriptors are **internally consistent** by three independent tests: (1) cross-validation stability: 1.07 ± 0.02 px RMS across 5 folds, with descriptor spread ± 0.4 mm (baseline), ± 0.15 mm (WD); (2) bootstrap 95 % confidence intervals: $b \in [23.48, 25.04]$ mm, $WD \in [64.64, 65.13]$ mm, $\theta \in [21.04^\circ, 22.72^\circ]$; (3) frame-count stability: descriptors vary by < 0.2 mm (baseline) and < 0.1 mm (WD) from 3 to 10 frames (Table 4).

External validation remains open. Confirming the absolute accuracy of the geometric descriptors requires either (a) manufacturer specifications for the specific microscope and optical configuration used to collect the Pycaso calibration images, or (b) an independent 3D measurement of the sub-pupil positions and working distance. Both are beyond the scope of this study and remain a limitation (Section 5.2).

4.6 Comparison with Generic Methods

We did not benchmark against Schöps et al. [24] or Pan et al. [17] on this dataset because their code targets pinhole-based initialisation and central projection with high-parameter-count spline or implicit-distortion corrections. Applying these methods to a non-central CMO microscope would require adapting their initialisation strategy to handle two sub-pupils, which is beyond the scope of this comparison. We expect that a generic spline model with sufficient parameters could match the Zernike rayfield’s 0.47 px RMS, but at the cost of thousands of non-interpretable parameters, compared to our 26 interpretable parameters.

4.7 Internal Validation

Three internal validation experiments assess the robustness of the 26p pipeline on the Pycaso dataset.

Cross-validation. 5-fold leave-2-out cross-validation (fit on 8 frames, evaluate on 2 held-out frames) yields a mean pixel RMS of 1.07 ± 0.02 px. The low standard deviation (~ 0.02 px) confirms that the

model is not overfitting to specific frames; the calibration is stable across the 10-frame dataset. Fold 1 shows slightly higher RMS (1.09 px) and a lower baseline estimate ($b = 23.1$ mm) because the held-out frames (0 and 1) cover the most extreme Z values.

Bootstrap descriptor stability. 100 bootstrap iterations (resampling 10 frames with replacement, refitting the full pipeline each time) give 95 % confidence intervals:

- Baseline b : [24.2, 25.4] mm (mean 24.6 mm, $\sigma = 0.4$ mm)
- Working distance WD : [64.7, 65.1] mm (mean 64.9 mm, $\sigma = 0.14$ mm)
- Convergence angle θ : [21.4°, 23.0°] (mean 22.2°, $\sigma = 0.4^\circ$)

The working distance is the most stable descriptor ($\sigma < 0.2$ mm), while the baseline and convergence angle show 0.4 mm and 0.4° variation respectively—consistent with the gauge sensitivity of the Zernike origin field noted in Section 5.2.

Sensitivity to fixed focal length f_x . Varying the fixed pinhole focal length f_x by $\pm 10\%$ (5 values: 23040, 24320, 25600, 26880, 28160 px) produces:

- Pixel RMS: $1.21 \rightarrow 0.95$ px—remains near the one-pixel range, confirming that the rayfield is stable as a calibration object.
- Working distance: $58.3 \rightarrow 71.3$ mm—linearly proportional to the initial f_x ($\pm 10\% f_x \rightarrow \pm 10\% WD$). WD is directly scaled by the pinhole reference and is *not* independently identifiable.
- Convergence angle: $24.1^\circ \rightarrow 20.5^\circ$ —varies inversely with WD .
- Baseline: $23.9 \rightarrow 25.1$ mm ($\pm 2.5\%$)—moderately stable.

The absolute physical descriptors, especially WD and θ , remain sensitive to the chosen pinhole reference scale. Descriptor readout therefore requires either an independently known optical scale (pixel pitch, magnification) or external metrological validation—a known limitation of all single-view calibration methods without a known reference object.

4.8 Application Sanity Check: Specimen Reconstruction

A non-metrological sanity check validates the calibrated rayfields on an independent speckled coin specimen imaged with the same Pycaso microscope (single stereo pair, no ChArUco pattern). This is **not** absolute metrological validation—no ground-truth 3D data exists for the coin—but it confirms that the rayfields produce coherent dense reconstruction on data unseen during calibration. Dense left-to-right pixel correspondences are computed with DIS optical flow over a 1448×1448 px ROI (92.7 % of pixels produce valid correspondences). For each correspondence, the two rays from the calibrated model are intersected at their closest point (ray-gap triangulation).

The median ray gap is 0.001 mm for both models—excellent coherence, confirming that the two channels' rays intersect each other at the sub-millimetre level, despite never having been jointly optimised on this specimen. The CMO 26p produces a smoother reconstructed surface (Z roughness, measured as median absolute deviation, of 0.072 mm vs. 0.176 mm for the Zernike), because the compact 26-parameter model imposes implicit spatial smoothness via its rigid-sub-pupil and affine-direction parameterisation, whereas the 57-parameter Zernike model has the flexibility to absorb detection noise into high-frequency




Figure 9. Dense specimen reconstruction. Top row: left image, horizontal disparity field, valid mask (92.7 % coverage). Bottom row: reconstructed Z (depth) maps from the CMO 26p and Zernike 57p rayfields, and the ray-gap histogram. Both models produce nearly identical ray gaps (median 0.001 mm), confirming rayfield coherence. The CMO 26p surface is smoother (Z roughness 0.072 mm) than the Zernike (0.176 mm) because the compact model imposes implicit smoothness via its parameterisation.

surface variations. OpenCV stereo reconstruction was not available because the standard OpenCV calibration failed to converge on this microscope’s calibration data.




Figure 10. Rigid-gauge removal: Zernike vs CMO surface comparison. Top row: CMO and Zernike surfaces after plane-normalization (relief only).

DeltaZ raw and plane-normalized residuals. Bottom row:

DeltaZ before SE(3) alignment (dominated by an affine ramp, $R^2 = 0.995$), the fitted affine plane, and the residual histogram. After SE(3)+plane removal, the median 3D residual drops from 3.8 mm to 0.06 mm — a 98.5 % reduction. The interpretation is Case A: the Zernike and CMO reconstructions differ primarily by a global reference-frame gauge, not by a genuine non-rigid rayfield discrepancy.

While the SE(3) alignment shows that both models reconstruct the same surface up to a rigid frame change, a residual depth-scale discrepancy remains: the Zernike reconstruction exhibits approximately 40 % larger relief amplitude than the CMO 26p model (standard deviation ratio $1.40\times$, Figure 10). To diagnose whether this represents a true scale difference or a depth-only effect, we fit three nested transformations from Zernike to CMO points on the matched specimen correspondences (50 000 subsampled points). An SE(3)-only fit (no scale) yields a median residual of 0.009 mm, already confirming excellent geometric agreement. Adding a global isotropic scale (Sim(3)) produces an optimal scale factor $s = 1.00$, ruling out a uniform 3D magnification—the difference is not a simple focal-length or pixel-scale error. However, an anisotropic scale model with independent (s_x, s_y, s_z) recovers $s_x \approx s_y \approx 1.00$ but $s_z = 0.67$, indicating a **depth-only scale difference** of approximately 33 %. The lateral dimensions (x, y) are identical between the two reconstructions; only the reconstructed depth range differs.

This depth-only scale discrepancy is consistent with a difference in effective stereo convergence or baseline between the flexible 57-parameter Zernike rayfield and the compact 26-parameter CMO model. The CMO model’s rigid sub-pupil and telecentric direction parameterisation impose a specific stereo geometry that the Zernike model, being a free rayfield, is not obligated to reproduce exactly. Since no independent 3D ground truth is available for the coin specimen, we interpret the depth-scale difference as a convergence-depth ambiguity inherent to uncalibrated stereo reconstruction, not as a failure of either model. The depth-only scale discrepancy is therefore reported as an unresolved convergence-depth ambiguity between the flexible Zernike rayfield and the compact physical model. It does not by itself establish which reconstruction is metrically closer to the real coin surface. Resolving this point would require an independent 3D reference, for example profilometry of the same specimen.

Table 5 quantifies the five reconstructions. The rayfield-initialised CMO model (before BA) has a median ray gap of 0.022 mm—an order of magnitude worse than all BA-refined variants (~ 0.001 mm). This was previously masked by a Y-axis coordinate convention issue (see `core/conventions.py`); with the corrected world frame, the initial model’s triangulation quality is revealed to be substantially worse than the BA-refined models. All BA variants recover tight ray intersections (median gap ~ 1 μ m), confirming that the ray geometry is correct regardless of the regularisation strategy. The regularised variants achieve Z MAD 0.027 mm, marginally better than the unregularised BA (0.030 mm)—the prior does not degrade the geometric reconstruction. All CMO BA variants converge to a magnification ratio of ~ 0.193 vs. the Zernike reference at 0.197, reflecting a 3.5 % systematic model bias between the compact 26-parameter physical model and the flexible 57-parameter rayfield.

Table 5. Five-variant specimen reconstruction: surface roughness (Z MAD), median ray gap, and magnification ratio vs. the nominal 18.75 mm 2-cent euro coin.

Model	Z MAD (mm)	Median gap (mm)	Magnification
Zernike rayfield (57 p)	0.194	0.0011	0.1968
CMO 26 p (rayfield init)	0.073	0.0224	0.1904
CMO 26 p — BA unregularised	0.030	0.0011	0.1931
CMO 26 p — BA + isotropic prior ($\alpha=10^{-2}$)	0.027	0.0011	0.1930
CMO 26 p — BA + Schur prior ($\alpha=10^{-3}$)	0.027	0.0011	0.1930




Figure 11. Five-variant specimen reconstruction with Schur-regularised variants. Each row shows the Z-map (surface relief after mean-plane subtraction), the XY footprint (coloured by relief), and the ray-gap histogram (log scale, shared range). All CMO variants are aligned to the Zernike reference frame via `core/conventions.py`. The rayfield initialisation (row 2) has a median gap 0.022 mm—20× larger than the BA-refined rows—revealed by the corrected world-frame Y axis. The regularised variants (rows 4–5) achieve Z MAD 0.027 mm, marginally better than the unregularised BA (0.030 mm) and 2.6× smoother than the Zernike reference (0.194 mm). All BA variants converge to magnification ~ 0.193 vs. Zernike 0.197 (3.5% systematic model bias).

4.9 Model Performance

Table 6. Model comparison on Pycaso CMO data. The 26-parameter CMO+SE(3) model is the compact physical reference.

Model	p	Ray RMS (mm)	Px RMS (px)	P50 (px)	P95 (px)
OpenCV stereo	—	—	> 300	—	—
Perspective CMO	19	3.48	~86	—	—
Telecentric CMO + shear	14	0.12	14.6	13.2	22.4
CMO + SE(3)	26	0.0021	1.06	0.87	1.84
CMO + SE(3) + corner BA	26	~0.0019	~0.98	~0.80	~1.70
Zernike O(0)+d(2)	57	0.0007	0.47	0.34	0.86

The 26-parameter CMO+SE(3) model achieves 1.06 px reprojection with less than half the parameters of the Zernike reference. It is fully interpretable: sub-pupils, focal length, telecentric slopes, and SE(3) arm corrections.




Figure 12. Two-panel BIC bar chart. (a) Ray-space BIC: the CMO telecentric family achieves the lowest BIC by a decisive margin (> 40 000 points). (b) Operational usability score with a 1.5 px reprojection guard: models exceeding the threshold are marked as REJECTED (shown in grey/orange); only the 26p SE(3)-aligned model passes the guard (shown in green).

4.10 BIC Model Selection

Table 7. BIC model selection on the Zernike rayfield. Models above 1.5 px reprojection are REJECTED by the operational guard.

Model	p	RMS (mm)	BIC_{ray}	S_{usable}	Status
cmo telecentric shear	14	0.111	-36 128	+978 890	REJECTED
cmo telecentric	12	0.146	-33 201	+986 044	REJECTED
pinhole parallel plate	6	6.337	5 079		
central pinhole	0	8.343	6 500		
central brown conrady	10	8.324	6 549		
CMO + SE(3) 26p	26	0.002	-32 433	-32 433	BEST USABLE

The ray-space BIC decisively selects the CMO telecentric family ($\Delta\text{BIC} > 40\,000$ over all non-CMO alternatives). The operational usability score correctly rejects the 14p base model as unusable and selects the 26p aligned model.

4.11 Geometric Descriptors

Table 8. CMO-consistent geometric descriptors read directly from the Zernike rayfield at the centre pixel.

Descriptor	Symbol	Value
Stereo baseline	b	24.9 mm
Sub-pupil depth	z_p	2.5 mm
Working distance	WD	64.7 mm
Objective focal length	f_{obj}	62.2 mm
Convergence angle	θ	22.6°

4.12 SE(3) Ablation

Table 9. SE(3) arm alignment ablation. Per-arm DOFs are individually necessary.

Variant	Params	Ray RMS (mm)	Px RMS (px)	P50 (px)	P95 (px)
Telecentric baseline	14	0.048	14.6	13.2	22.4
+ Rotation only	20	0.014	3.74	2.68	7.13
+ Translation only	20	0.010	2.44	2.06	4.15
+ Full SE(3)	26	0.0021	1.06	0.87	1.84
+ Shared rotation	23	0.0083	—	—	—
+ Shared translation	23	0.0041	—	—	—
+ Differential only	20	0.0070	4.04	2.85	7.74

Both rotation and translation are essential; neither alone suffices. Per-arm DOFs are individually necessary—any compression degrades the fit (+92 % to +289 % ray RMS increase).

5 Discussion

5.1 The Rayfield as Diagnostic Instrument

The key methodological contribution is the feedback loop: Ray2D preprocessing $\rightarrow$ Ray3D measurement $\rightarrow$ residual modal analysis $\rightarrow$ hypothesis $\rightarrow$ model refinement $\rightarrow$ verification. Without the rayfield, the 2D reprojection error tells you *that* the model is wrong but not *how*. The Zernike projection of Δd and Δm directly identifies whether the missing DOF is global (low-order Zernike modes) or spatial (high-order modes).

Figure 6 illustrates this diagnostic value concretely. The residual heatmap at each model stage is a **visual hypothesis generator**: the perspective CMO’s structured d_y error (23.6° RMS) sends the user to telecentric models; the telecentric model’s uniform Z_0^0 piston (0.27° RMS) sends the user to global arm alignment; and the SE(3)-refined model’s near-zero residual (0.003° RMS) confirms convergence. At no point does the user need to guess—the residual pattern at each stage dictates the next hypothesis.

In this case study, the Z_0^0 -dominated residual correctly guided us to the SE(3) arm alignment, while alternative hypotheses (pre-warp, variable origin) were correctly rejected by the same diagnostic.

5.2 Limitations

1. **Single dataset.** Results are demonstrated on 10 stereo pairs from one Pycaso CMO microscope. Generalisation to other CMO designs, manufacturers, or magnification settings has not been demonstrated and requires additional validation.
2. **No independent 3D ground truth.** All residuals (ray-space and pixel-equivalent) are computed on the same ChArUco board points used for calibration. There is no validation against a separate certified 3D object or an independently measured reference.
3. **Effective rayfield descriptors are gauge-sensitive.** The geometric descriptors (b , WD , f_{obj} , θ) are read from the rayfield under the transverse gauge $O \cdot d = 0$ and a fixed pinhole reference $f_x = 25600$. The working distance is linearly proportional to the initial f_x ($\pm 10\% f_x \rightarrow \pm 10\% WD$, Section 4.7), and is *not* independently identifiable without an external scale reference.
4. **SE(3) translation parameters are not uniquely identifiable.** The rotation angles ($\sim 2.5^\circ$, $\sim 3.7^\circ$) are stable across optimisation runs, but the translation components trade off with the telecentric base parameters (WD , f_{obj} , b , principal point). Only the SE(3) rotation is a robust diagnostic; the translation is effective rather than absolute.
5. **The Zernike rayfield is a training residual.** Its 0.0007 mm ray-space RMS is a self-evaluation on the support points it was fitted to; the 0.47 px pixel RMS is a noise-floor estimate, not an independent measurement. The pixel RMS column in all tables provides the apples-to-apples comparison.
6. **Fixed pinhole reference.** The Zernike BA uses a fixed $f_x = 25600$ and principal point at image centre. Errors in this reference propagate to the rayfield estimates and the derived physical descriptors.
7. **Absolute metrological validation would require:** (a) manufacturer specifications for the actual microscope used, (b) an independent measurement of WD or baseline, or (c) a certified 3D calibration object at a known distance. None of these are available for the open Pycaso dataset.

5.3 Generalisability

The methodology—measure with a flexible non-parametric basis, read physical descriptors, diagnose residual structure, build compact physical model, validate by BIC—is transferable to any non-standard optical instrument (endoscopes, Scheimpflug cameras, multi-camera arrays). The SE(3) arm alignment is specific to CMO microscopes, but the residual analysis that identified it is general.

5.4 Why the Two-Stage Decomposition Works

The central insight of this work is not a specific model architecture but a problem decomposition strategy. Direct bundle adjustment jointly estimates optical parameters θ and board poses η from 2D corners; for non-central optics, the Schur complement of the reduced camera system [32] reveals a coupling norm of 0.81, meaning optics and poses are nearly inseparable—the problem is severely ill-conditioned. By inserting the Zernike rayfield as an intermediate variable, the estimation splits into two well-conditioned sub-problems: (1) fit a flexible rayfield to corners (Zernike BA, coupling norm

0.0—poses are eliminated by the constrained-pose model), and (2) identify physical optics in ray space (pose-free model selection with BIC).

This decomposition is the mechanism behind all three empirical findings: the TPS pre-processing stabilises step (1) by cleaning the 2D input; the rayfield decouples step (2) from pose estimation; and the rayfield-initialised parameters lie so close to the corner optimum that direct BA can barely improve them.

6 Conclusion

We have demonstrated the first real-data validation of the StereoComplex non-central calibration pipeline on a CMO stereo microscope. The key results are:

- A compact 26-parameter physical model (quasi-telecentric CMO skeleton + per-channel SE(3) arm alignment) achieves 1.06 px reprojection (P50=0.87 px, P95=1.84 px) where standard OpenCV stereo calibration fails (> 300 px).
- The Zernike rayfield serves as a diagnostic instrument: residual modal analysis identifies missing degrees of freedom, and the operational usability score with reprojection guard correctly selects the usable model.
- Effective rayfield descriptors—baseline, working distance, focal length, convergence angle—are read directly from the measured rays (within the chosen rayfield gauge and reference scale), without fitting a physical model.
- The rayfield-initialised parameters are near-optimal for direct corner bundle adjustment (17 % improvement).

The CMO Pycaso case study validates a generalisable methodology: measure first, model second, diagnose by residual, refine by hypothesis testing. Code, data, and documentation are publicly available at <https://github.com/jeffwitz/StereoComplex>. A Zenodo archive of the submitted version will be deposited before journal submission.

Prospective

Several directions extend this work. First, the methodology should be validated on additional CMO microscope designs and manufacturers to confirm generalisability beyond the single Pycaso instrument. Second, the Greenough stereo microscope family—with tilted optical axes rather than parallel sub-pupils—presents a harder identification problem where the telecentric model does not apply; the residual-guided construction method would need to discover a different geometric skeleton. Third, quantitative uncertainty propagation from ChArUco corner noise to physical descriptor estimates (baseline, working distance) would strengthen the metrological claim. Fourth, deep learning-based corner detection could replace the hand-tuned Hessian+TPS pipeline, potentially improving the raw 2D input quality and further reducing the gauge instability. Finally, extending the Zernike rayfield to time-varying deformations (e.g., vibration, thermal drift) would enable *in-situ* calibration monitoring without dedicated targets.

Reproducibility Statement

The StereoComplex codebase (commit 60272b7, branch develop) and all artefacts required to reproduce these results are available at <https://github.com/jeffwitz/StereoComplex>. The raw

Pycaso calibration images are available from <https://github.com/LaboratoireMecaniqueLille/Pycaso>. To reproduce the numerical results without access to the raw images, restart from the pre-computed intermediate state (`intermediate_state.npz`) which contains the already-detected, Hessian-completed, and TPS-denoised corner positions, the fitted Zernike rayfield, and the initial 26p model parameters. A Zenodo archive of the exact submitted version will be deposited before journal submission; until then, the GitHub commit hash should be used as the reproducibility identifier.

Author Contributions

J.-F. Witz conceived the study, developed the StereoComplex framework, implemented all methods, conducted the experiments, and wrote the manuscript.

Acknowledgments

The author thanks Claude (Anthropic) and ChatGPT (OpenAI) for assistance with code generation, manuscript drafting, and critical review of the documentation.

Code and Data Availability

- StereoComplex: <https://github.com/jeffwitz/StereoComplex> (Code: GPL-2.0-or-later. Documentation: CC BY-SA 4.0.)
- Pycaso dataset: <https://github.com/LaboratoireMecaniqueLille/Pycaso>
- All generated artefacts: `docs/assets/pycaso_real_data/`
- Intermediate state (restart without raw images): `docs/assets/pycaso_real_data/intermediate_state.npz`

References

- [1] H. Akaike. A new look at the statistical model identification. *IEEE Transactions on Automatic Control*, 19(6):716–723, 1974. doi:[10.1109/tac.1974.1100705](https://doi.org/10.1109/tac.1974.1100705).
- [2] A. Bartoli and P. Sturm. Structure-from-motion using lines: Representation, triangulation, and bundle adjustment. *Computer Vision and Image Understanding*, 100(3):416–441, 2005. doi:[10.1016/j.cviu.2005.06.001](https://doi.org/10.1016/j.cviu.2005.06.001).
- [3] F. L. Bookstein. Principal warps: Thin-plate splines and the decomposition of deformations. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 11(6):567–585, 1989. doi:[10.1109/34.24792](https://doi.org/10.1109/34.24792).
- [4] J.-Y. Bouguet. Camera calibration toolbox for MATLAB, 2008. URL: http://www.vision.caltech.edu/bouguetj/calib_doc/.
- [5] D. C. Brown. Decentering distortion of lenses. *Photogrammetric Engineering*, 32(3):444–462, 1966. doi:[10.14358/pers.32.3.444](https://doi.org/10.14358/pers.32.3.444).
- [6] F. Camposeco, T. Sattler, and M. Pollefeys. Non-parametric camera model calibration with splines. In *Proceedings of the German Conference on Pattern Recognition (GCPR)*, pages 285–296, 2016. doi:[10.1007/978-3-319-45886-1_23](https://doi.org/10.1007/978-3-319-45886-1_23).
- [7] E. Caron, J.-F. Witz, C. Cuvier, A. Beaurain, V. Magnier, and A. El Bartali. PYCASO: Python module for calibration of cameras by Soloff’s method. *SoftwareX*, 23:101440, 2023. doi:[10.1016/j.softx.2023.101440](https://doi.org/10.1016/j.softx.2023.101440).
- [8] S. Garrido-Jurado, R. Muñoz Salinas, F. J. Madrid-Cuevas, and M. J. Marín-Jiménez. Automatic generation and detection of highly reliable fiducial markers under occlusion. *Pattern Recognition*, 47(6):2280–2292, 2014. doi:[10.1016/j.patcog.2014.01.005](https://doi.org/10.1016/j.patcog.2014.01.005).
- [9] M. D. Grossberg and S. K. Nayar. A general imaging model and a method for finding its parameters. In *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, pages 108–115, 2001. doi:[10.1109/iccv.2001.937611](https://doi.org/10.1109/iccv.2001.937611).
- [10] S. Inoué and K. R. Spring. *Video Microscopy: The Fundamentals*. Plenum Press, New York, 2nd edition, 1997. doi:[10.1007/978-1-4757-6925-8](https://doi.org/10.1007/978-1-4757-6925-8).
- [11] J. Kannala and S. S. Brandt. A generic camera model and calibration method for conventional, wide-angle, and fish-eye lenses. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 28(8):1335–1340, 2006. doi:[10.1109/tpami.2006.153](https://doi.org/10.1109/tpami.2006.153).
- [12] Z. Liu, J. Guo, W. Shi, X. Chen, and H. Xie. Generic distortion model for metrology under optical microscopes. *Optics and Lasers in Engineering*, 103:119–126, 2018. doi:[10.1016/j.optlaseng.2017.12.006](https://doi.org/10.1016/j.optlaseng.2017.12.006).
- [13] C. Mei and P. Rives. Single view point omnidirectional camera calibration from planar grids. In *Proceedings of the IEEE International Conference on Robotics and Automation (ICRA)*, pages 3945–3950, 2007. doi:[10.1109/robot.2007.364084](https://doi.org/10.1109/robot.2007.364084).

- [14] P. Miraldo and H. Araujo. Calibration of smooth camera models. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 35(9):2091–2103, 2013. doi:[10.1109/TPAMI.2012.258](https://doi.org/10.1109/TPAMI.2012.258).
- [15] R. J. Noll. Zernike polynomials and atmospheric turbulence. *Journal of the Optical Society of America*, 66(3):207–211, 1976. doi:[10.1364/josa.66.000207](https://doi.org/10.1364/josa.66.000207).
- [16] OpenCV. Open source computer vision library. URL: <https://opencv.org>.
- [17] L. Pan, M. Pollefeys, and V. Larsson. Camera pose estimation using implicit distortion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 12819–12828, 2022. doi:[10.1109/CVPR52688.2022.01248](https://doi.org/10.1109/CVPR52688.2022.01248).
- [18] J. Plücker. On a new geometry of space. *Philosophical Transactions of the Royal Society of London*, 155:725–791, 1865.
- [19] S. Ramalingam and P. Sturm. A unifying model for camera calibration. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 39(7):1309–1319, 2017. doi:[10.1109/tpami.2016.2592904](https://doi.org/10.1109/tpami.2016.2592904).
- [20] S. Ramalingam, P. Sturm, and S. K. Nayar. A generic structure-from-motion framework. In *Proceedings of the European Conference on Computer Vision (ECCV)*, pages 218–231, 2010. doi:[10.1007/978-3-642-15552-9_16](https://doi.org/10.1007/978-3-642-15552-9_16).
- [21] F. J. Romero-Ramirez, R. Muñoz Salinas, and R. Medina-Carnicer. Speeded up detection of squared fiducial markers. *Image and Vision Computing*, 76:38–47, 2018. doi:[10.1016/j.imavis.2018.05.004](https://doi.org/10.1016/j.imavis.2018.05.004).
- [22] J. Sanderson. *Understanding Light Microscopy*. Wiley, 2019. doi:[10.1002/9781118696736](https://doi.org/10.1002/9781118696736).
- [23] D. Scaramuzza, A. Martinelli, and R. Siegwart. A toolbox for easily calibrating omnidirectional cameras. In *Proceedings of the IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 5695–5701, 2006. doi:[10.1109/iros.2006.282372](https://doi.org/10.1109/iros.2006.282372).
- [24] T. Schöps, V. Larsson, M. Pollefeys, and T. Sattler. Why having 10,000 parameters in your camera model is better than twelve. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 2535–2544, 2020. doi:[10.1109/cvpr42600.2020.00261](https://doi.org/10.1109/cvpr42600.2020.00261).
- [25] H. W. Schreier, D. Garcia, and M. A. Sutton. Advances in light microscope stereo vision. *Experimental Mechanics*, 44(3):278–288, 2004. doi:[10.1177/0014485104041546](https://doi.org/10.1177/0014485104041546).
- [26] G. Schwarz. Estimating the dimension of a model. *Annals of Statistics*, 6(2):461–464, 1978. doi:[10.1214/aos/1176344136](https://doi.org/10.1214/aos/1176344136).
- [27] S. M. Soloff, R. J. Adrian, and Z. C. Liu. Distortion compensation for generalized stereoscopic particle image velocimetry. *Measurement Science and Technology*, 8(12):1441–1454, 1997. doi:[10.1088/0957-0233/8/12/008](https://doi.org/10.1088/0957-0233/8/12/008).
- [28] C. Strecha et al. On benchmarking camera calibration and multi-view stereo for high resolution imagery. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 1–8, 2008. doi:[10.1109/cvpr.2008.4587706](https://doi.org/10.1109/cvpr.2008.4587706).

- [29] P. Sturm and S. Ramalingam. A generic concept for camera calibration. In *Proceedings of the European Conference on Computer Vision (ECCV)*, pages 1–13, 2004. doi:[10.1007/978-3-540-24671-8_1](https://doi.org/10.1007/978-3-540-24671-8_1).
- [30] B. Triggs, P. McLauchlan, R. Hartley, and A. Fitzgibbon. Bundle adjustment — a modern synthesis. In *International Workshop on Vision Algorithms*, pages 298–372, 1999. doi:[10.1007/3-540-44480-7_21](https://doi.org/10.1007/3-540-44480-7_21).
- [31] S. Weber, N. Demmel, and D. Cremers. Power bundle adjustment for large-scale 3D reconstruction. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 281–289, 2023. doi:[10.1109/CVPR52729.2023.00035](https://doi.org/10.1109/CVPR52729.2023.00035).
- [32] J.-F. Witz. StereoComplex: Non-central stereo calibration framework, 2025. URL: <https://github.com/jeffwitz/StereoComplex>.
- [33] F. Zernike. Diffraction theory of the knife-edge test and its improved form, the phase-contrast method. *Monthly Notices of the Royal Astronomical Society*, 94:377–384, 1934. doi:[10.1093/mnras/94.5.377](https://doi.org/10.1093/mnras/94.5.377).
- [34] D. Zhang, M. Luo, and D. D. Arola. Displacement/strain measurements using an optical microscope and digital image correlation. *Optical Engineering*, 45(3):033605, 2006. doi:[10.3788/iop57.021506](https://doi.org/10.3788/iop57.021506).
- [35] Z. Zhang. A flexible new technique for camera calibration. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 22(11):1330–1334, 2000. doi:[10.1109/34.888718](https://doi.org/10.1109/34.888718).